Model Evaluation For Humans

About me

Husband and Father
Philly data scientist
Wrote a haiku once

Disclaimers

- Most useful for new(er) data scientists
- But there may be take aways for everyone
- Oversimplication
- Guidelines/examples are not exhaustive
- Breadth not depth
- Sometimes the obvious has to be said aloud

Goal of this talk

Discuss guidelines and practices for robust model evaluation

Outline

- Motivation
- Before development
  - Business & Data Understanding
- During development
  - Data Preparation → Modeling → Evaluation
- Deployment and beyond
  - Deployment → Evaluation & Monitoring → Iteration

Motivation: all models are wrong

"All models are wrong, ..."

$$\underset{\theta}{\operatorname{argmin}}\, L(y, \hat{y})$$

$$MLE := \underset{\theta}{\operatorname{argmax}}\, \pi(D \mid \theta)$$

$$\hat{y} = f(x) + \epsilon$$

$$P(\theta \mid D) \propto P(D \mid \theta) P(\theta)$$

Ask me how I know...

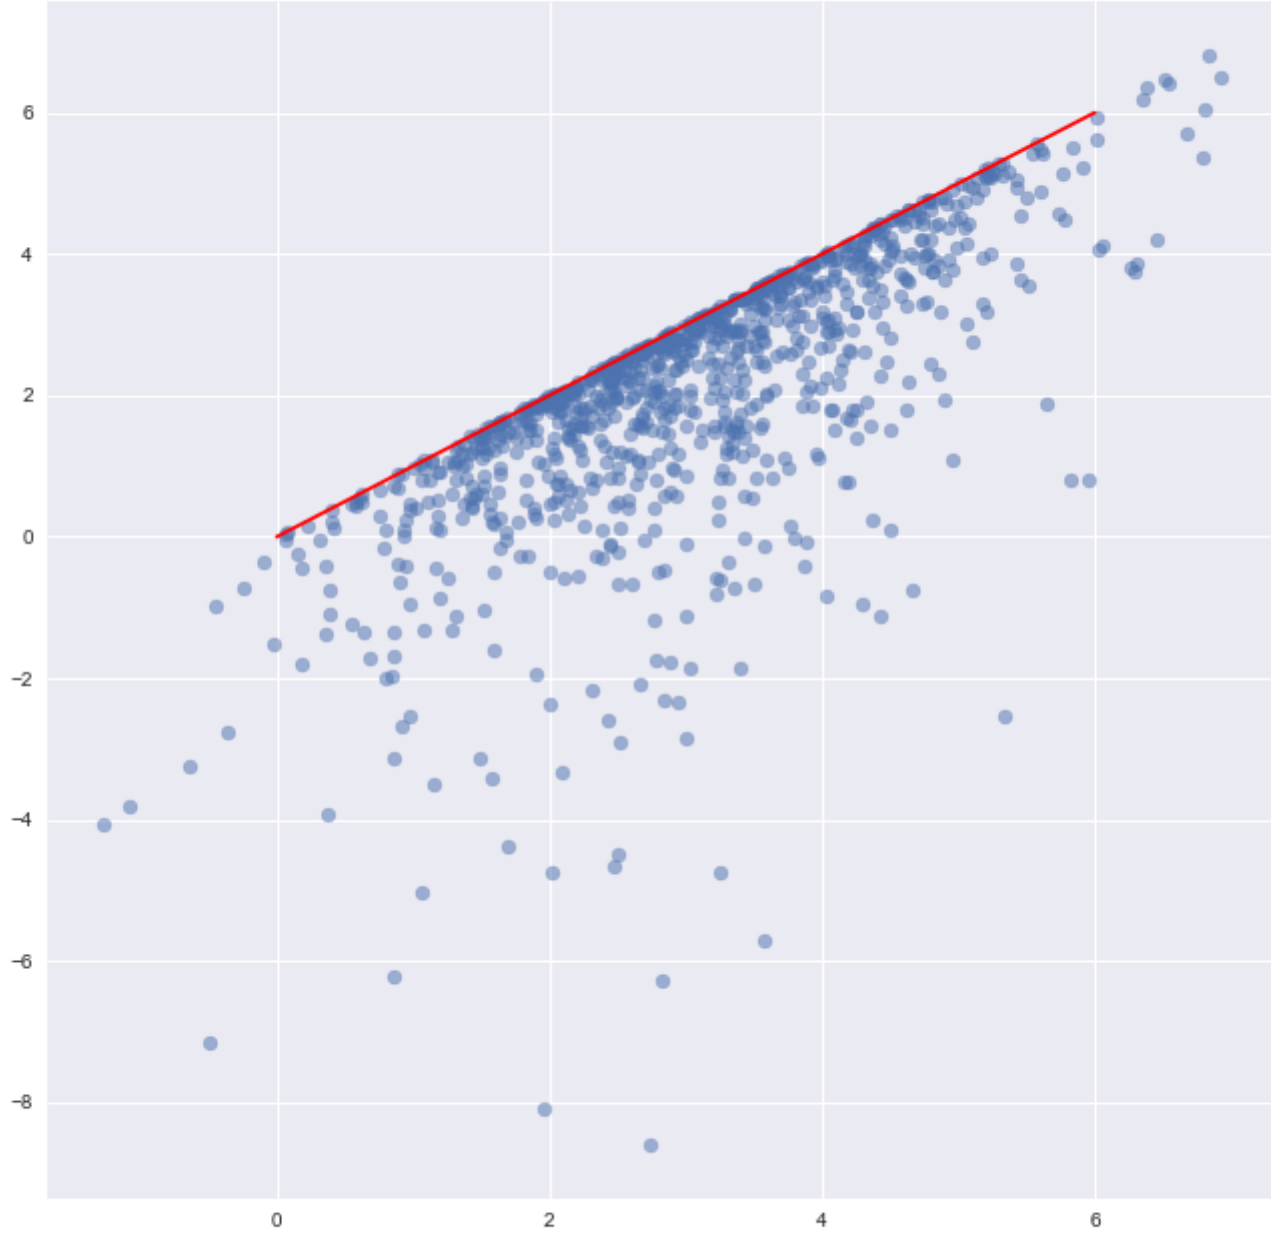
6
4
2
0
−2
−4
−6
−8
0
2
4
6

On a better day... still wrong

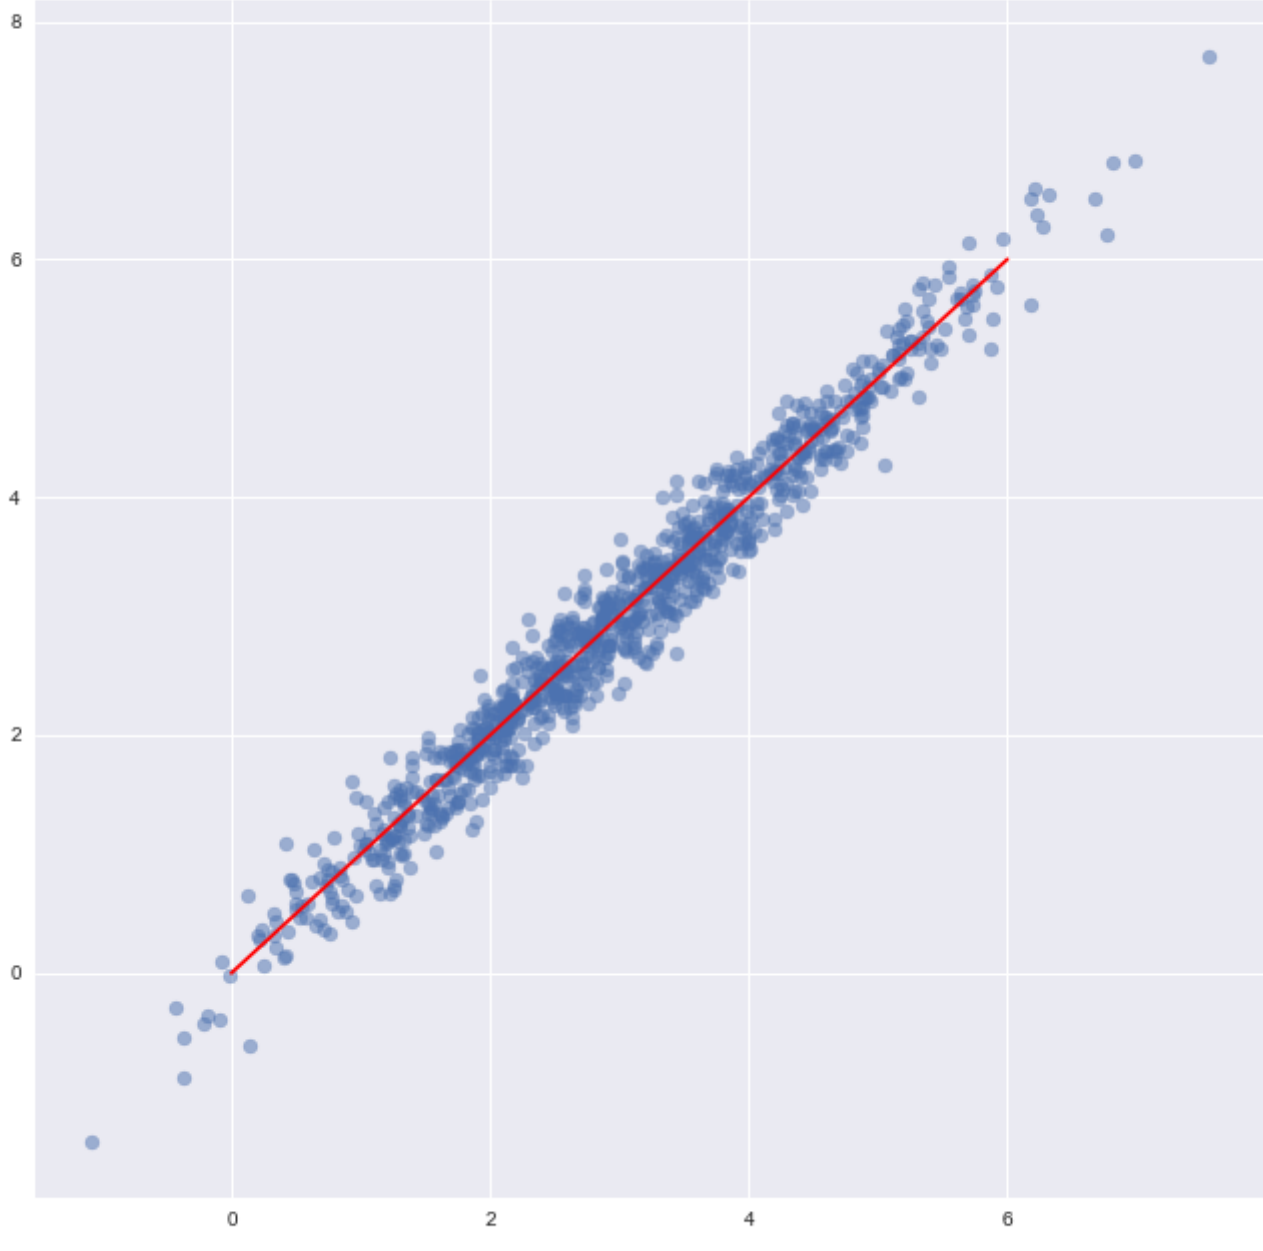
8
6
4
2
0
0
2
4
6

The zen of model evaluation

"*All models are wrong, some are useful*"

The zen of model evaluation summarized

- All models are wrong
- Some are useful
- Others are not
- Usefulness counts

Before development

- Highly context specific
- Essential for rest of project

Hazards of getting this wrong

- Model does not solve POI
- Model indirectly solves POI
- Model cannot be evaluated in prod
  - Reduced confidence in model
- Diminishing returns
- Iteration is impaired

Some things to think about...

Think about framing the problem

- How will this model be evaluated?
- How will success be *measured*?
- Is this a classification or regression problem?
- Choose metrics and KPIs
  - Conversion to business value
  - Comparability
    - Across time, data, iterations, etc.

Think about deployment and future iterations

- "Nothing is more permanent than temporary"
  - Prepare for deployment
  - Prepare for future iterations

Think about the feedback loop

- How will outcomes from the model be collected?
- Example: Ad recommendations
- Example: Intervention for hospital readmission
  - Is the model that bad or the intervention that good?

During development

- Most generalizable
- Opportunity adopt and standardize best practices

Test assumptions

- $f(\quad) =$
- Check distributions of data
- Plot, plot, plot
- Check features against conventional wisdom
- Etc.

Testing assumptions: Example

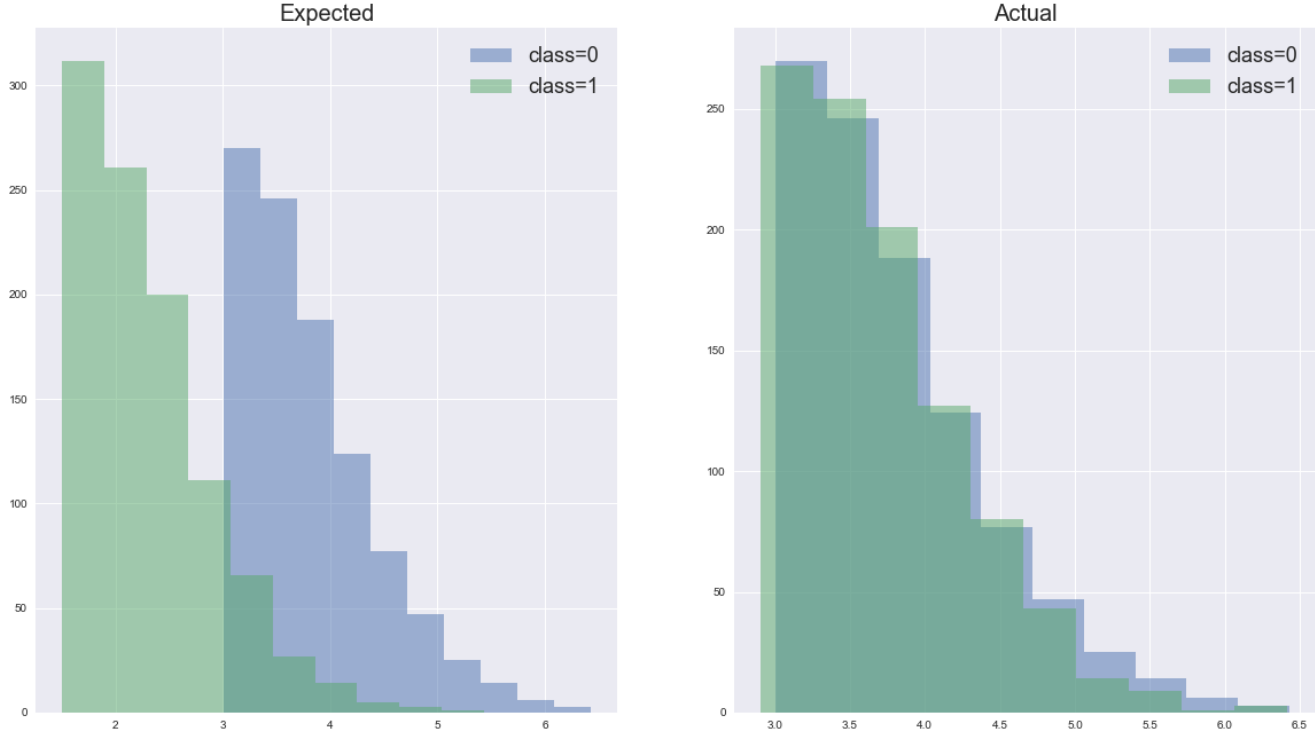

Test assumptions: Example

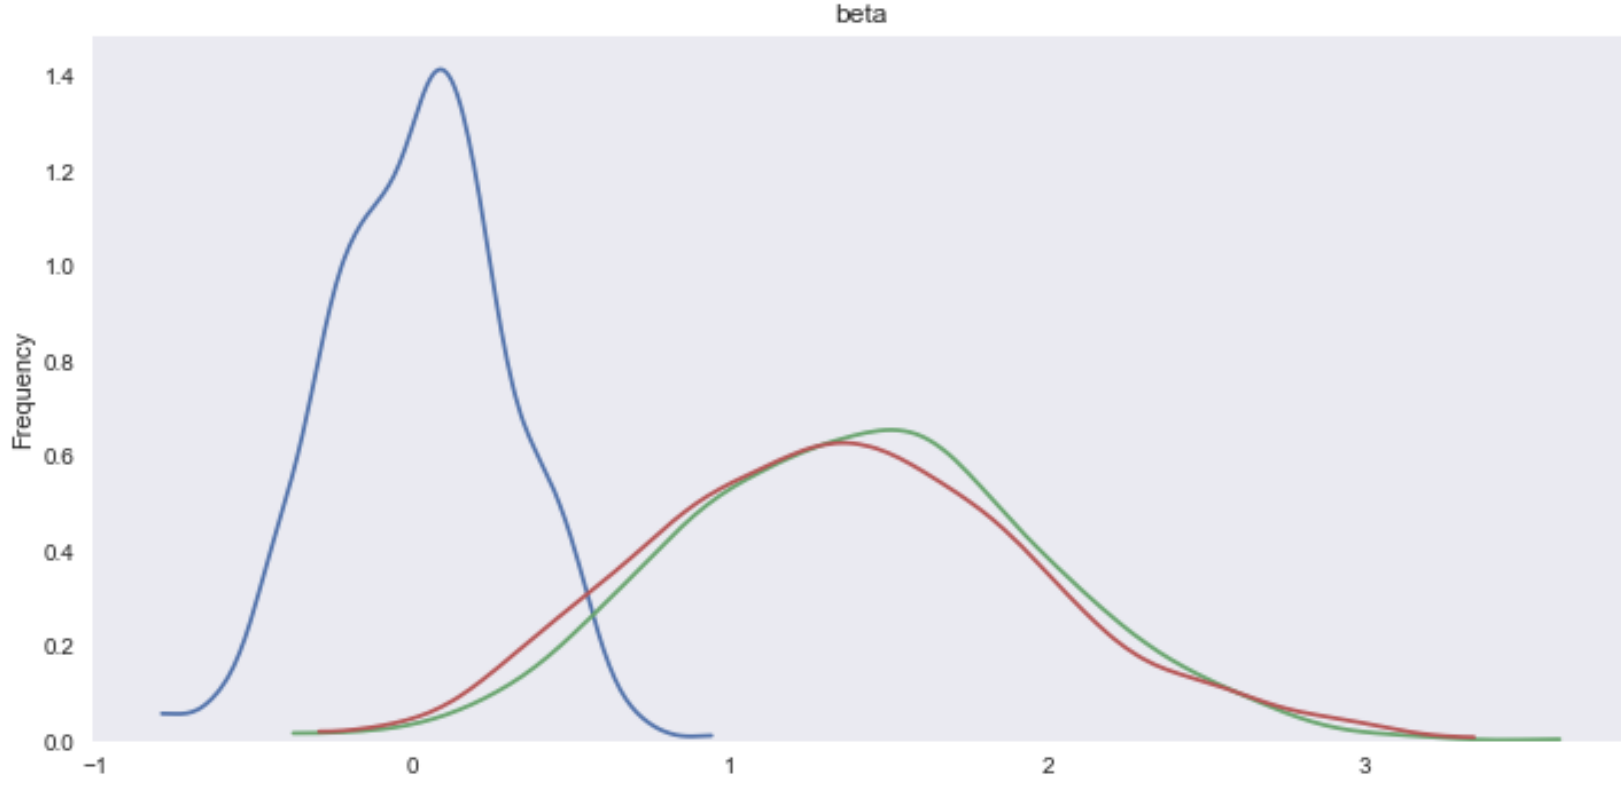

Establish a baseline

A good baseline is

- Easy to implement
- Easy to understand

Examples of baselines

- A simple heuristic
  - Example: Persistency
- An interpretable model
  - Example: Comorbidities
- The existing model or methodology
- Acceptance criteria

Cross validation

Replicate the deployment pattern

- Avoid leakage from future examples
- More like backtesting, less like K-fold CV
  - Train on Jan. predict on Feb.
  - Train on Feb. predict on Mar.

Train on a cleaner data set

- Validate modeling approach apart from noise in data set
  - Example: Ratings

Deployment and beyond

- This is what really matters (Remember the Zen!)
- Ensure that your model is (and stays) useful!
- Iterate with confidence!

Why A/B testing?

Monitoring:

- Keep it simple and actionable
  - Example: Favor Range over KL Divergence

- Reuse sanity checks from development

Conclusion

- Hopefully you have some take aways
- Share your suggestions!
- https://dantegates.github.io/2019/01/07/model-evaluation-for-humans.html (https://dantegates.github.io/2019/01/07/model-evaluation-for-humans.html)

Questions?